

Allocator with Implicit Free List and Coalescing

Current allocator works... performance?

Utilization


Can create too much fragmentation

Throughput


Can take a long time to find a block

Choosing a Free Block


`p2 = malloc(1)`



Choosing a Free Block



`p2 = malloc(1)`




`p4 = malloc(4)`

First fit: Use (and possibly split) the first block that works

danger of fragmentation

Choosing a Free Block


`p2 = malloc(1)`



Choosing a Free Block

`p2 = malloc(1)`

`p4 = malloc(4)`



Best fit: Use (and possibly split) the smallest block that works
usually reduces fragmentation

First-Fit Implementation

Our current implementation is first-fit:

```
while (GET_SIZE(HDRP(bp)) != 0) {  
    if (!GET_ALLOC(HDRP(bp))  
        && (GET_SIZE(HDRP(bp)) >= new_size)) {  
        set_allocated(bp, new_size);  
        return bp;  
    }  
    bp = NEXT_BLK(P(bp));  
}
```

Best-Fit Implementation

A best-fit search:

```
void *best_bp = NULL;

while (GET_SIZE(HDRP(bp)) != 0) {
    if (!GET_ALLOC(HDRP(bp))
        && (GET_SIZE(HDRP(bp)) >= new_size)) {
        if (!best_bp
            || (GET_SIZE(HDRP(bp)) < GET_SIZE(HDRP(best_bp))))
            best_bp = bp;
    }
    bp = NEXT_BLKPTR(bp);
}


if (best_bp) {
    set_allocated(best_bp, new_size);
    return best_bp;
}
```

[Copy](#)

Trades throughput for utilization


Internal Fragmentation

```
p = mm_malloc(8);  
memset(p, 0, 8);
```




Internal Fragmentation

```
p = mm_malloc(8);  
memset(p, 0, 8);
```



Internal Fragmentation

```
p = mm_malloc(8);  
memset(p, 0, 8);
```




Everything except the application payload reduces utilization

Internal fragmentation refers to space within an allocated block that is unusable to the application

Internal Fragmentation


```
p = mm_malloc(8);  
memset(p, 0, 8);
```



Sources of internal fragmentation:

- headers and footers
- empty space to maintain alignment
- empty space due to choice of fit


Encoding Header and Footer Information



```
typedef struct {  
    size_t size;  
    char    allocated;  
} block_header;
```

```
typedef struct {  
    size_t size;  
    int filler;  
} block_footer;
```

Encoding Header and Footer Information




```
typedef size_t block_header;
```

```
typedef size_t block_footer;
```

Since a block size is always a multiple of 16, low 4 bits are always 0

Idea: use the low bit to indicate allocation status

Encoding Header and Footer Information



```
typedef size_t block_header;
```

```
typedef size_t block_footer;
```

```
#define GET(p) (*(size_t *) (p))
```

```
#define GET_ALLOC(p) (GET(p) & 0x1)
```

```
#define GET_SIZE(p) (GET(p) & ~0xF)
```

```
#define PUT(p, val) (*(size_t *) (p) = (val))
```

```
#define PACK(size, alloc) ((size) | (alloc))
```

Packing Demo

```
#include <stdio.h>
#include <stdlib.h>

#define GET(p)      (*(size_t *) (p))
#define PUT(p, val) (*(size_t *) (p) = (val))

#define GET_ALLOC(p) (GET(p) & 0x1)
#define GET_SIZE(p)  (GET(p) & ~0xF)

#define PACK(size, alloc) ((size) | (alloc))


int main() {
    void *p = malloc(sizeof(size_t));

    PUT(p, PACK(48, 1));
    printf("%ld %s\n",
           GET_SIZE(p),
           (GET_ALLOC(p) ? "alloc" : "unalloc"));
}
```

[Copy](#)


Alignment

A smaller header can break our alignment strategy:



Solution:

- Make sure **first_bp** has correct alignment




- Align total block size, not payload size

Even Smaller Headers and Footers

If the block size is constrained to be $< 2^{32}$:


```
typedef int block_header;  
typedef int block_footer;  
#define GET(p)      (*(int *) (p))  
#define PUT(p, val) (*(int *) (p) = (val))
```




Allocator might treat very large blocks differently

Advanced: Footers Only for Unallocated Blocks

Our allocator needs to go backwards only for coalescing:



Idea: record in block header whether *previous* block is allocated




Make sure block size is big enough for footer to be added


free (p2) \Rightarrow previous block is unallocated, so use **PREV_BLK**

Advanced: Footers Only for Unallocated Blocks

Our allocator needs to go backwards only for coalescing:




Idea: record in block header whether *previous* block is allocated




Make sure block size is big enough for footer to be added

`free(p3)` \Rightarrow previous block is allocated, so don't try `PREV_BLKPTR`


Looking for an Unallocated Block



Looking for an Unallocated Block




Looking for an Unallocated Block



Finding an unallocated is a significant limitation on throughput

Looking for an Unallocated Block




Finding an unallocated is a significant limitation on throughput

Instead of searching through *all* blocks, keep a list of just the free ones


The allocator will have an **explicit free list** instead of an **implicit free list**

Explicit Free Lists

Make sure that every block has room for a pointer
replaces the application's payload




free (p3)




Explicit Free Lists

Make sure that every block has room for a pointer
replaces the application's payload




`free (p3)`




Explicit Free Lists


Make sure that every block has room for a pointer
replaces the application's payload




free (p3)



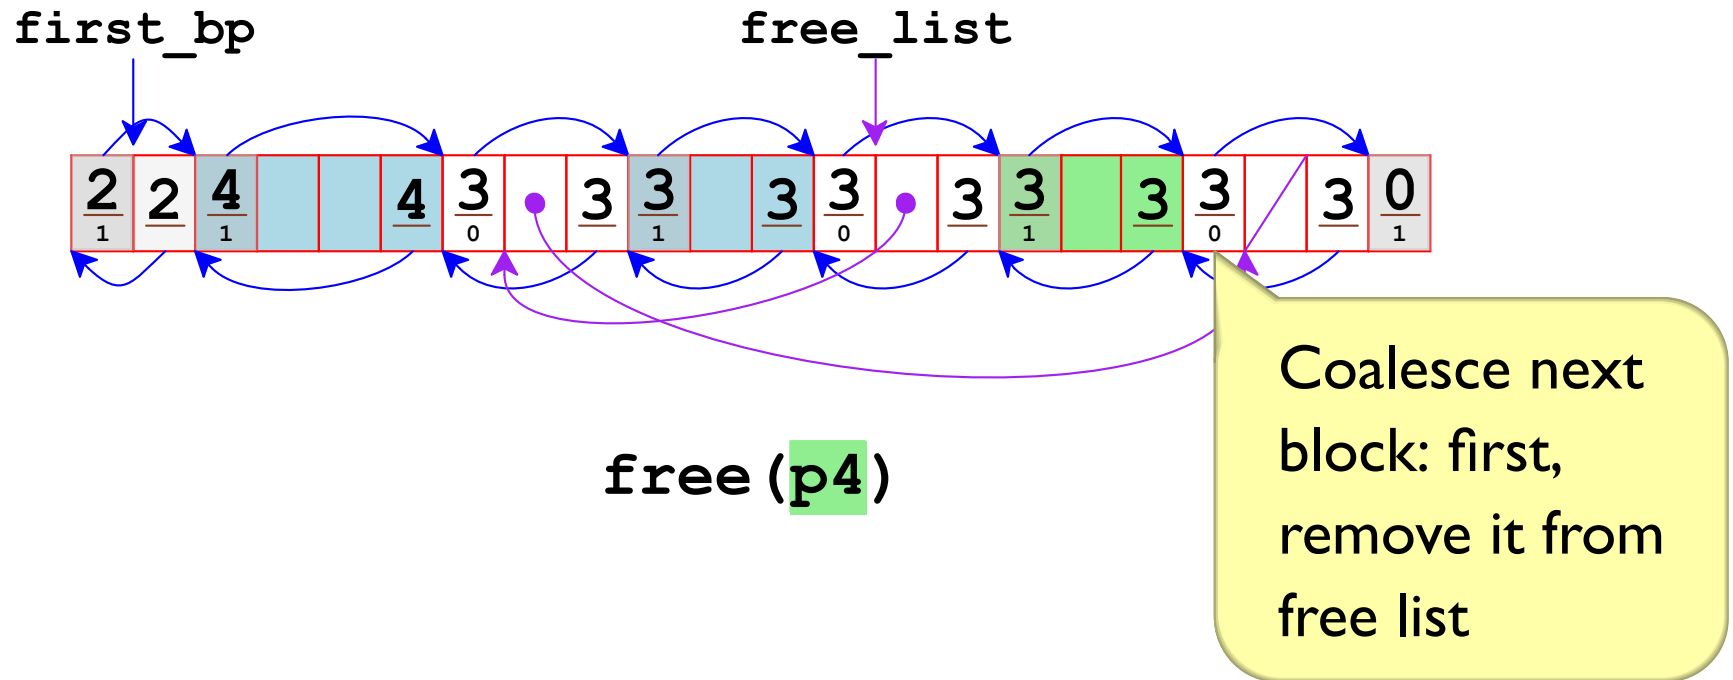
Free Lists and Coalescing




Free Lists and Coalescing




Free Lists and Coalescing




Free Lists and Coalescing




free (p4)




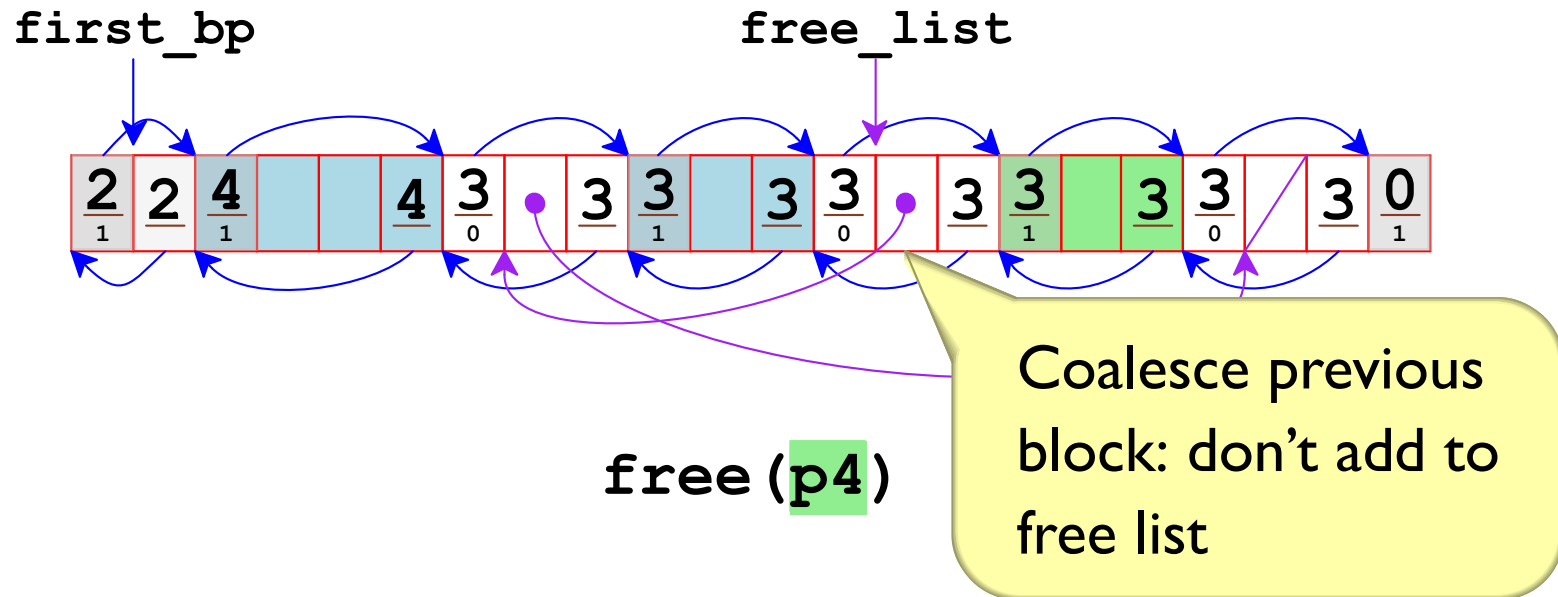
Free Lists and Coalescing




free (p4)




Free Lists and Coalescing



Free Lists and Coalescing



free (p4)



Invariant: every unallocated block is on the free list

Free List Data Structure

Linked list not convenient?

Use a doubly-linked lists

```
typedef struct list_node {
    struct list_node *prev;
    struct list_node *next;
} list_node;

....

void *mm_malloc(size_t size) {
    int need_size = max(size, sizeof(list_node));
    int new_size = ALIGN(need_size + OVERHEAD);
    ....
}
```

Free List Data Structure

Linked list not convenient?

Use a doubly-linked lists

```
void *coalesce(void *bp) {  
    ....  
  
    if (prev_alloc && next_alloc) {      /* Case 1 */  
        add_to_free_list((list_node *)bp);  
    }  
  
    ....  
}
```

Selecting from a Free List

First fit and **best fit** are still options with an explicit free list

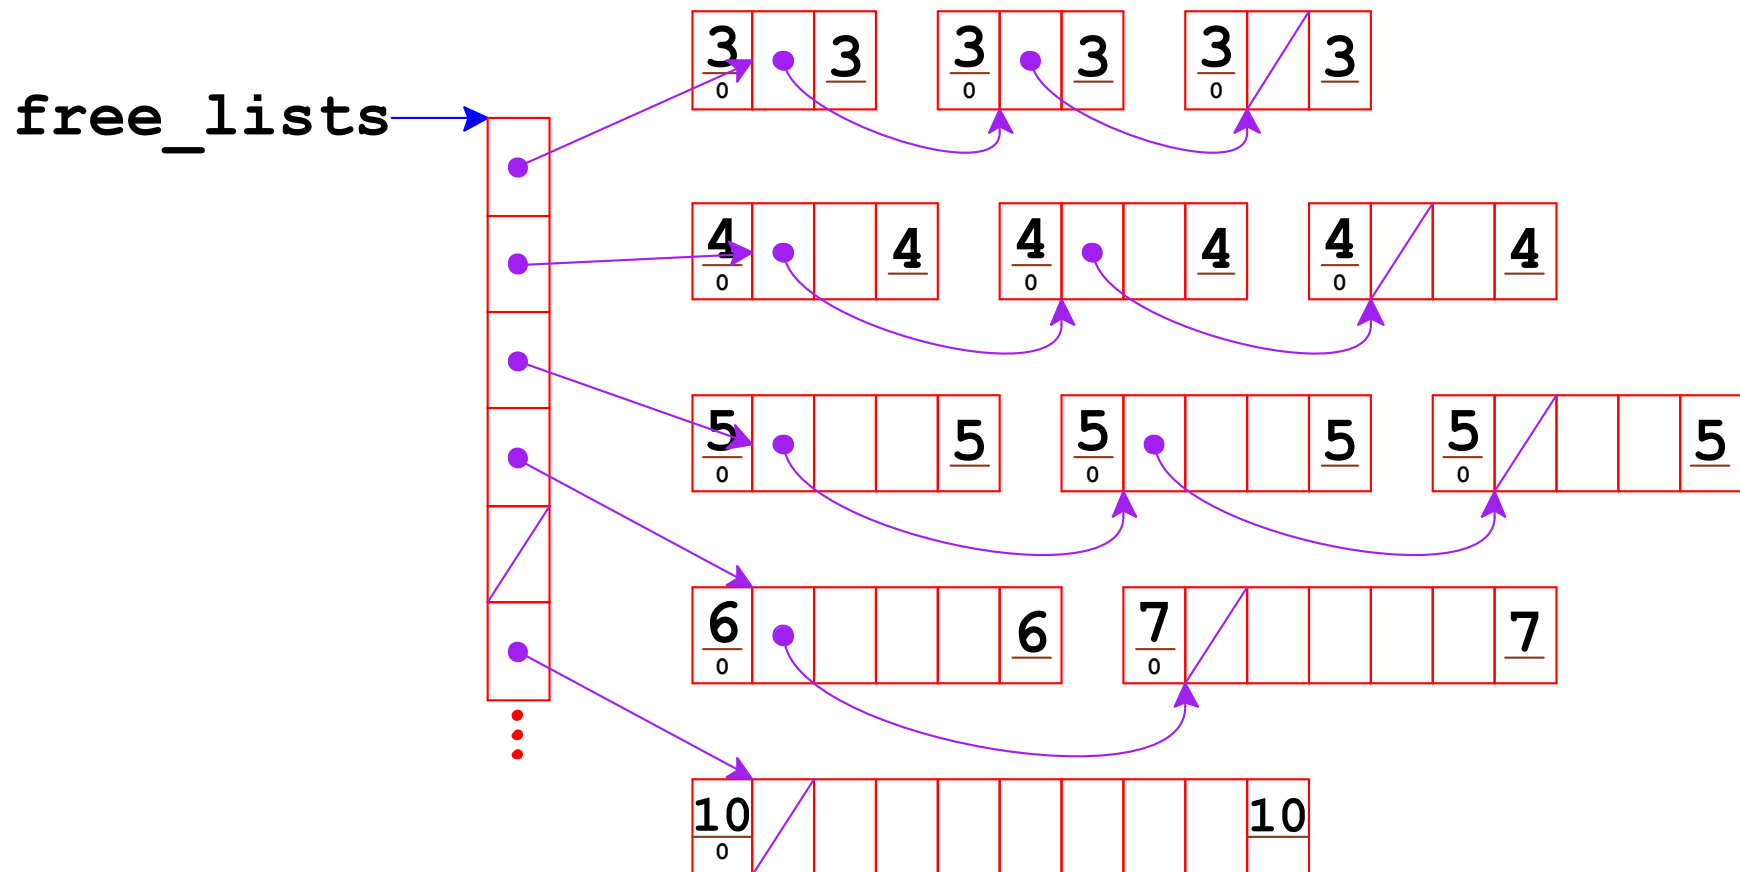
More options:

LIFO — add to front of free list; take from front of list
a kind of first-fit that tends to promote locality

address ordered — pick block with lowest address
may reduce fragmentation

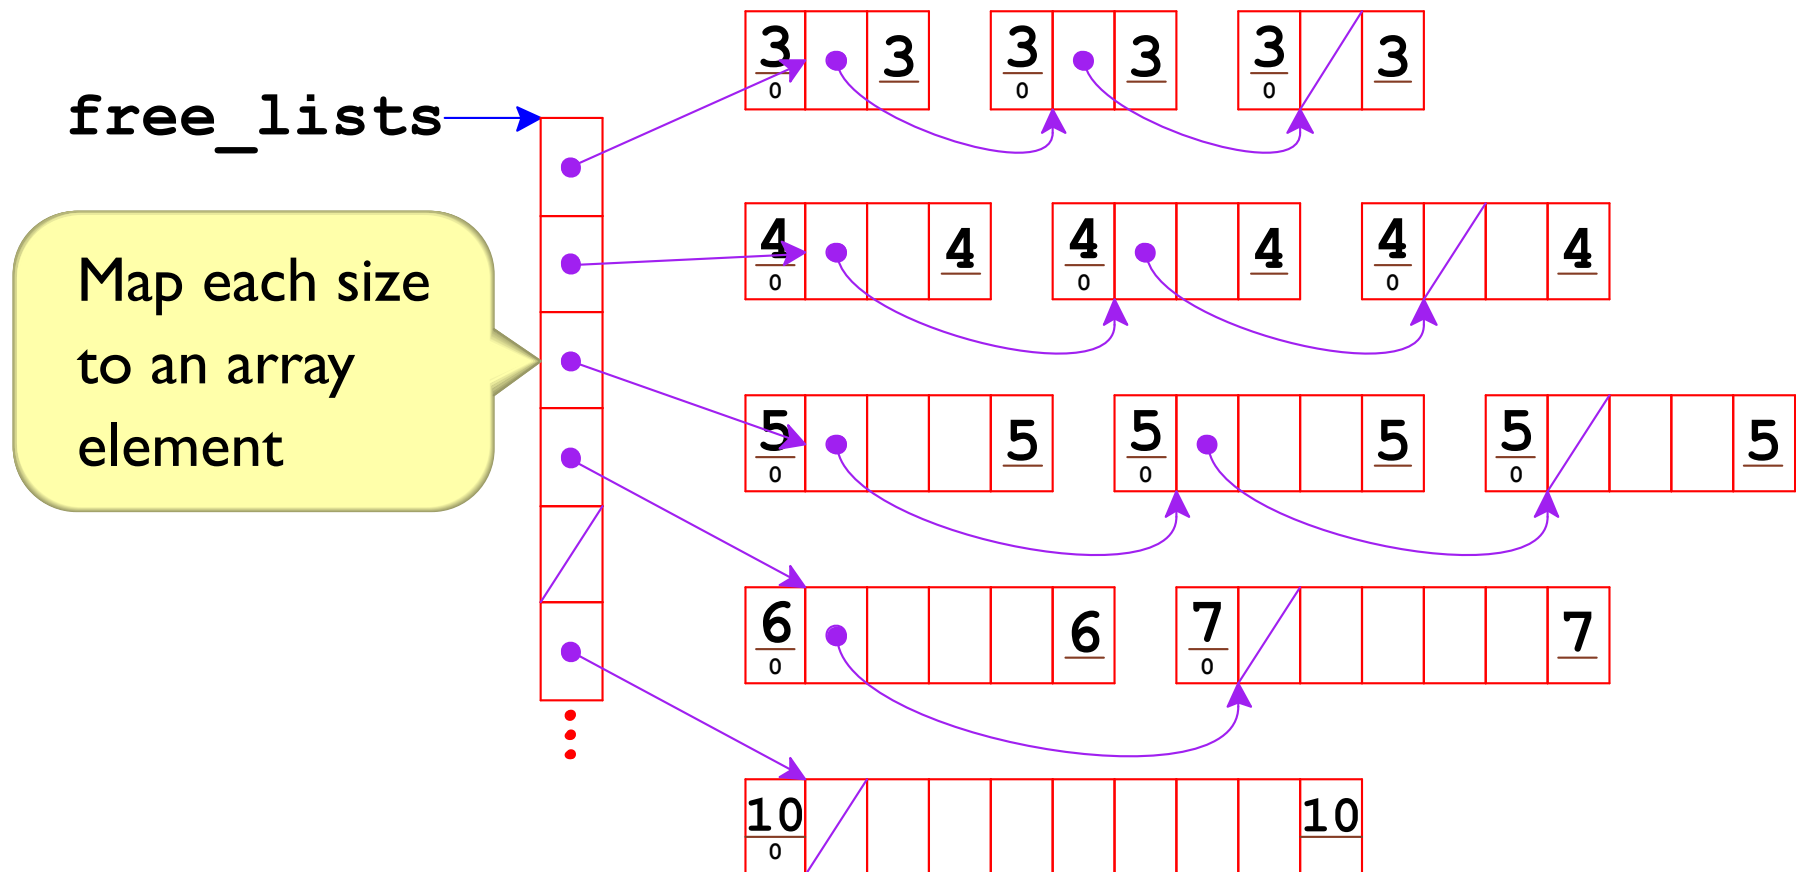
Segregated Free List

A **segregated free list** is an array of free lists, where each list has objects of a particular size class



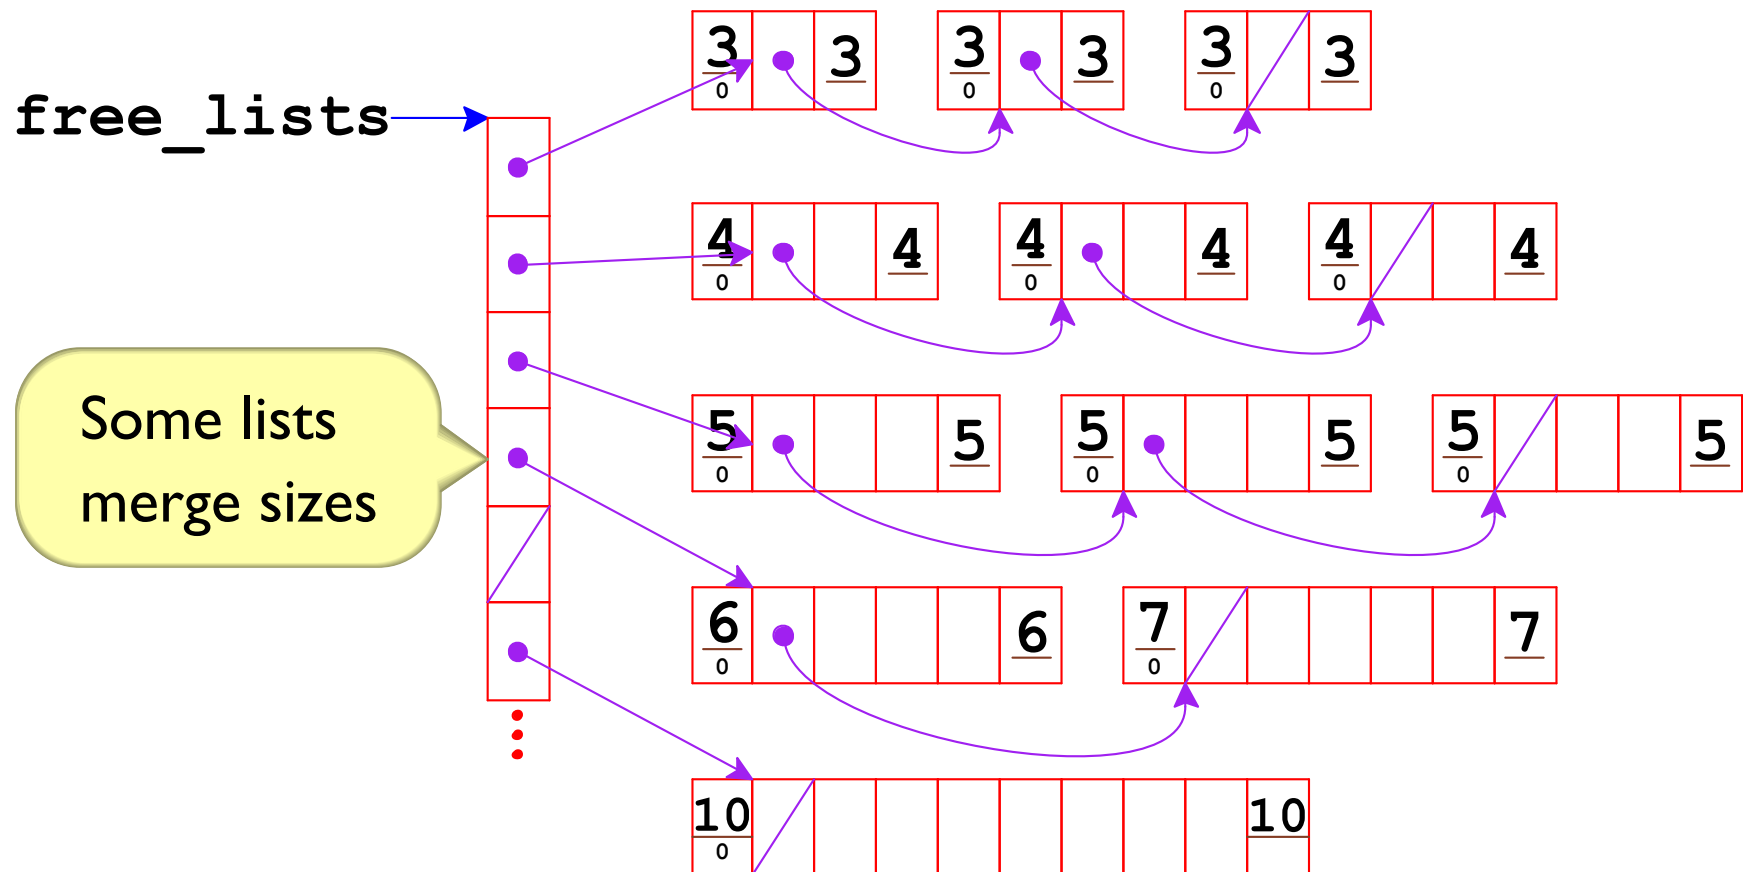
Segregated Free List

A **segregated free list** is an array of free lists, where each list has objects of a particular size class




Segregated Free List

A **segregated free list** is an array of free lists, where each list has objects of a particular size class




Balanced Binary Tree

Instead of a free list, a **free tree** can support efficient best-fit



Building an Allocator with `mmap`

Our allocator implementation so far depends on a contiguous heap:



Allocators can use `mmap` instead of `sbrk`


using `mmap` works in more environments

Unlike `sbrk`, separate `mmap` calls don't always return contiguous addresses

Building an Allocator with mmap

Chain together mapped pages as mini **sbrk**-like allocators

first_chunk




Building an Allocator with `mmap`

Chain together mapped pages as mini **sbrk**-like allocators

`first_`

In each chunk from `mmap`, use first few bytes for chaining




Building an Allocator with `mmap`

Chain together mapped pages as mini `sbrk`-like allocators

Can compute per-chunk `first_bp`
as offset from page address


`first_chur`



Building an Allocator with mmap

Chain together mapped pages as mini **sbrk**-like allocators


first_chunk




Building an Allocator with mmap

Chain together mapped pages as mini **sbrk**-like allocators

first_chunk




Can remove empty page from page list and use **munmap**



Building an Allocator with mmap

Chain together mapped pages as mini **sbrk**-like allocators

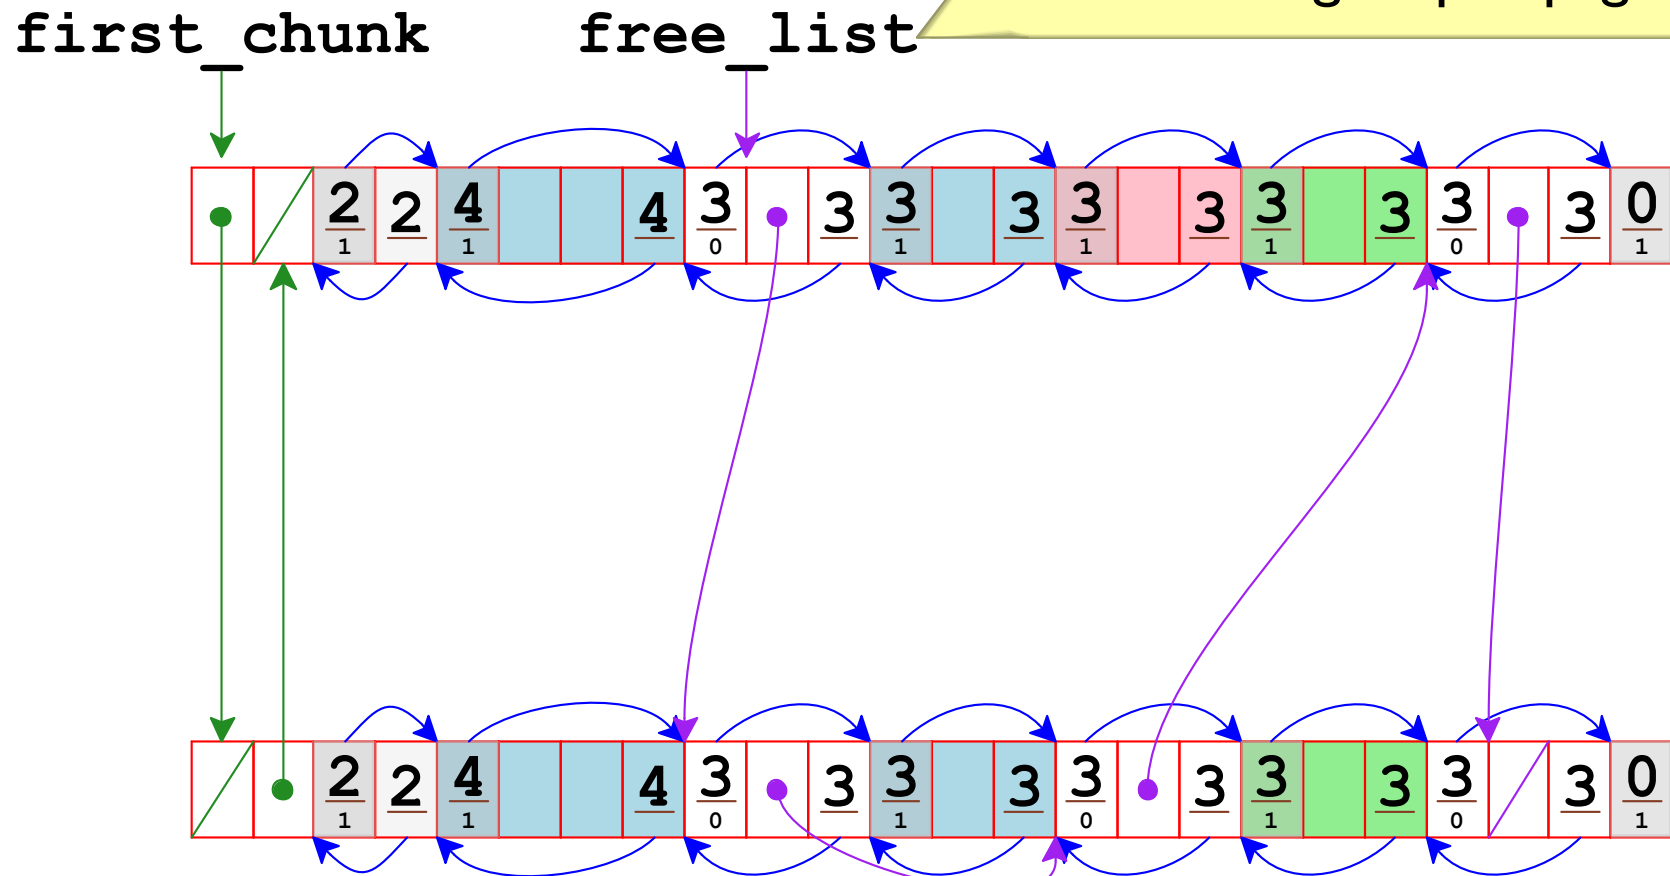
first_chunk



Building an Allocator with mmap

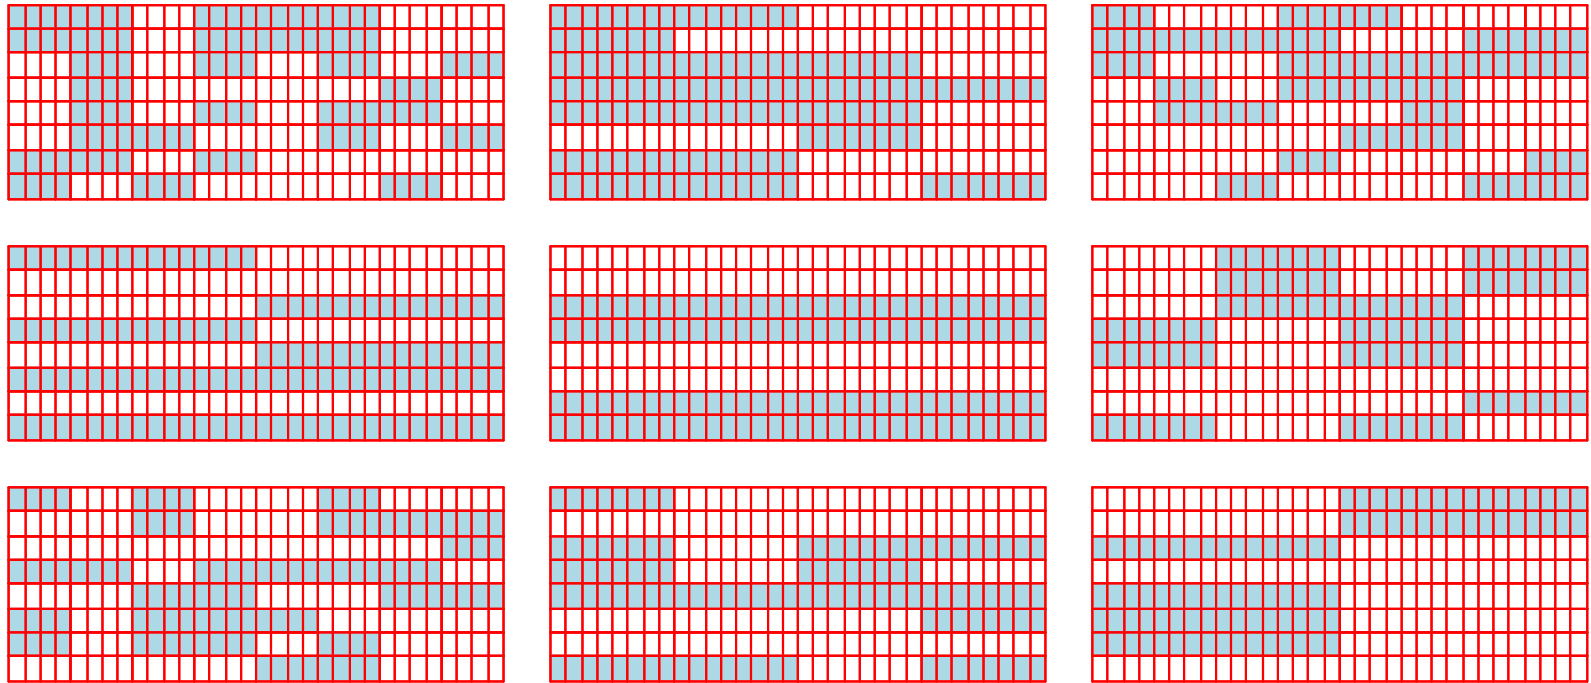
Chain together mapped pages as mini **sbrk**-like allocators

Free list might span pages



Something Completely Different: Segregated Allocation

Each chunk of memory has uniform-sized blocks




I How does **free** know an allocated block's size?

Based on the address: the allocator keeps a mapping of address ranges to block sizes

Something Completely Different: Segregated Allocation

Each chunk of memory has uniform-sized blocks




2 How is unallocated space represented?

Through a free list or chunk-specific bitmap

Something Completely Different: Segregated Allocation

Each chunk of memory has uniform-sized blocks




3 How is unallocated space selected for each allocation?

Any unallocated block will work within a chunk that holds the block size

Something Completely Different: Segregated Allocation

Each chunk of memory has uniform-sized blocks




4 How finely is unallocated space tracked?

Block sizes must be rounded up to match some chunk's content

Something Completely Different: Segregated Allocation

Each chunk of memory has uniform-sized blocks



5 When are more pages needed from the kernel?

When no chunk for the block size has any unallocated blocks